0.1 Congruence Path Refinement

In this section we present an algorithm that uses an existing congruence relation and refines it to the point where equivalent states can be "merged".

Definition 0.1.1. Let $\mathcal{A} = (Q, \Sigma, q_0, \delta, c)$ be a DPA and let $\equiv \subseteq Q \times Q$ be an equivalence relation on the state set. For every equivalence class $\kappa \subseteq Q$, let $r_\kappa \in \kappa$ be an arbitrary representative of that class. For a DPA $\mathcal{A}' = (Q', \Sigma, q_0', \delta', c')$, we say that $\mathcal{A}'$ is a *representative merge of $\mathcal{A}$ w.r.t. $\equiv$* if it satisfies the following:

- $Q' = \{r_{[q]_\equiv} \subseteq Q \mid q \in Q\}$
- $q_0' = r_{[q_0]_\equiv}$
- For all $q \in Q'$ and $a \in \Sigma$: $\delta'(q, a) = r_{[\delta(q,a)]_\equiv}$
- $c' = c\restriction_{Q'}$

Definition 0.1.2. Let $\mathcal{A} = (Q, \Sigma, q_0, \delta, c)$ be a DPA and let $R \subseteq Q \times Q$ be a congruence relation on the state space. For any states $p, q \in Q$, let $L_{[p]_R \to [q]_R} \subseteq \Sigma^*$ be the set of words w such that for any $u \sqsubseteq w$, $(\delta(p, u), q) \in R$ iff $u \in \{\varepsilon, w\}$. In other words, the set contains all minimal words by which the automaton reaches $[q]_R$ from $[p]_R$.

Let $\kappa \subseteq Q$ be an equivalence class of R and let $p, q \in \kappa$. We define $R_\kappa \subseteq \kappa \times \kappa$ as the smallest set $(p, q) \in R_\kappa$ iff the following holds for all words $w \in L_{\kappa \to \kappa}$:

- $(\delta^*(p, w), \delta^*(q, w)) \in R_\kappa$
- $\min\{c(\delta^*(p, u)) \mid u \sqsubset w\} = \min\{c(\delta^*(q, u)) \mid u \sqsubset w\}$

Finally, we call $\equiv^R_{\mathrm{PR}} = \bigcup_{q \in Q} R_{[q]_R}$ the *path refinement of R*.

Lemma 0.1.1. *The path refinement is a well defined equivalence relation.*

Proof. We have to consider the sets $L_{[p]_R \to [q]_R}$ and the sets R_κ. For $L_{[p]_R \to [q]_R}$, the definition works because R has the congruence property.

For R_κ, consider the following function $f : 2^{Q \times Q} \to 2^{Q \times Q}$:

$$f(X) = \{(p, q) \in Y_\kappa \mid \text{for all } w \in L_{\kappa \to \kappa} : (\delta^*(p, w), \delta^*(q, w)) \in X\}$$

$$Y_\kappa = \{(p, q) \in Q \times Q \mid \text{for all } w \in L_{\kappa \to \kappa} : \min\{c(\delta^*(p, u)) \mid u \sqsubset w\} = \min\{c(\delta^*(q, u)) \mid u \sqsubset w\}\}$$

Now Let $X_0 = \{(p, q) \in \kappa \times \kappa \mid p = q\}$ and $X_{i+1} = f(X_i)$. f is monotone w.r.t. $\subseteq$, so there must be a fixed point X_∞. By Kleene's fixed point theorem and from the definition of R_κ, we have $X_\infty = \mathrm{lfp}(f) = R_\kappa$.

Every X_i is an equivalence relation on κ: for $i = 0$, every state is only equivalent to itself, and for $i > 0$, the three properties can easily be verified via induction. Hence, $X_\infty = R_\kappa$ is also an equivalence relation. All R_κ are disjoint and thus $\equiv^R_{\mathrm{PR}}$ has to be an equivalence relation as well. □

Theorem 0.1.2. *Let $\mathcal{A} = (Q, \Sigma, q_0, \delta, c)$ be a DPA and let $R \subseteq Q \times Q$ be a congruence relation that implies language equivalence. Let $\mathcal{B}$ be a representative merge of $\mathcal{A}$ w.r.t. $\equiv^R_{\mathrm{PR}}$. Then $L(\mathcal{A}) = L(\mathcal{B})$.*

Proof. Let n be the number of non-trivial equivalence classes in $\equiv^R_{\mathrm{PR}}$, i.e. classes with size greater than 1. If $n = 0$, then $p \equiv^R_{\mathrm{PR}} q$ iff $p = q$ and therefore $\mathcal{B} = \mathcal{A}$.

Now assume for an argument of induction that the statement is true for n and we want to show that it still holds for $n + 1$ classes. Let $\kappa \subseteq Q$ be an arbitrary non-trivial equivalence class of $\equiv^R_{\mathrm{PR}}$. Let $\mathcal{A}' = (Q', \Sigma, q'_0, \delta', c')$ be the representative merge of $\mathcal{A}$ w.r.t. $\equiv^R_{\mathrm{PR}}\restriction_\kappa$ with the same representative r_κ as in $\mathcal{B}$. The path refinement equivalence of $\mathcal{A}'$ then is equal to $\equiv^R_{\mathrm{PR}}\restriction_{Q'}$ and has n non-trivial equivalence classes (as κ was merged into a single state). By induction, $L(\mathcal{A}') = L(\mathcal{B})$. It remains to be proven that $L(\mathcal{A}) = L(\mathcal{A}')$.

Let $\alpha \in \Sigma^\omega$ be a word with runs $\rho \in Q^\omega$ and $\rho' \in (Q')^\omega$ of $\mathcal{A}$ and $\mathcal{A}'$ respectively. Let $\lambda \subseteq Q$ be the equivalence class of R from which κ was extracted.

Claim 1: At every position i, $\rho(i) \in \kappa$ iff $\rho'(i) \in \kappa$.
Let k_0 be the first position at which $\rho(k_0) \in \kappa$ is true. For all $i < k_0$, we have $\rho(i) = \rho'(i)$, and at k_0 we have $\rho(k_0) \equiv^R_{\mathrm{PR}} r_\kappa = \rho'(k_0)$.

Now assume that the claim holds for all $i \leq k$, where k is a position at which $\rho(k) \in \kappa$. Let $l > k$ be the next position at which $\rho(l) \in \lambda$. If l does not exist, then neither $\rho(i)$ nor $\rho'(i)$ are elements of κ for any $i > k$.

Let $w = \alpha[k, l]$. Since $\kappa \subseteq \lambda$, $w \in L_{\lambda\to\lambda}$. By definition of $\equiv^R_{\mathrm{PR}}$, that means $\delta^*(\rho(k), w) = \rho(l) \equiv^R_{\mathrm{PR}} \delta^*(\rho'(k), w)$. Between k and $l-1$, no redirected edge is used in ρ', so $\delta^*(\rho'(k), \alpha[k, l-1]) = \rho'(l-1)$. Finally, $\rho'(l) = \delta'(\rho'(l-1), \alpha(l)) = r_{[\delta(\rho'(l-1),\alpha(l)]_{\equiv^R_{\mathrm{PR}}}} \equiv^R_{\mathrm{PR}} \delta(\rho'(l-1), \alpha(l)) = \delta^*(\rho'(k), w)$. Thus, $\rho(l) \equiv^R_{\mathrm{PR}} \rho'(l)$.

Now, if $\rho(l) \in \kappa$, then $\rho(l + 1) \in \kappa$ and our proof of induction is complete. If $\rho(l) \notin \kappa$, then $\rho'(l) = \rho(l)$, so the runs visit the same states in all positions until κ is reached again. This also completes the proof of our claim.

Claim 2: If κ only occurs finitely often in ρ and ρ', then ρ is accepting iff ρ' is accepting.

Let $k \in \mathbb{N}$ be the last position at which $\rho(k)$ and $\rho'(k)$ are in κ. From this point on, $\rho'[k, \omega]$ is also a valid run of $\mathcal{A}$ on $\alpha[k, \omega]$. $\rho(k) \equiv^R_{\mathrm{PR}} \rho'(k)$, so $(\rho(k), \rho'(k)) \in R$. As R implies language equivalence, reading $\alpha[k, \omega]$ from either state in $\mathcal{A}$ leads to the same acceptance status. This also means that $\rho'(k)$ has the same acceptance status as $\rho(k)$.

Claim 3: If κ occurs infinitely often in ρ and ρ', then ρ is accepting iff ρ' is accepting.

Let $(k_i)_{i\in\mathbb{N}}$ be all positions at which κ is visited. For each k_i, let $l_i > k_i$ be the minimal position at which $\rho(l_i) \in \lambda$. In two steps, we first show that $c(\rho[l_i, k_{i+1}]) = c'(\rho'[l_i, k_{i+1}])$ and second that $\min \mathrm{Occ}(c(\rho[k_i, l_i])) = \min \mathrm{Occ}(c'(\rho'[k_i, l_i]))$. Together, these results mean that the minimal priority that is seen infinitely often in the two runs is the same.

First, observe that at every l_i, we either have $l_i = k_{i+1}$ (if $\rho(l_i) \in \kappa$) or $\rho(l_i) = \rho'(l_i)$. In the first case, $\rho[l_i, k_{i+1}]$ is empty, so $c(\varepsilon) = c'(\varepsilon)$ is true. In the second case, $\rho[l_i, k_{i+1}] = \rho'[l_i, k_{i+1}]$ and therefore $c(\rho[l_i, k_{i+1}]) = c'(\rho'[l_i, k_{i+1}])$.

Second, let $w_i = \alpha[k_i, l_i]$. Then $\alpha \in L_{\lambda\to\lambda}$ and $\min \mathrm{Occ}(c(\rho[k_i, l_i])) = \min \mathrm{Occ}(c'(\rho'[k_i, l_i]))$ holds directly by definition of $\equiv^R_{\mathrm{PR}}$. □

0.1.1 Algorithmic Definition

The definition of path refinement that we introduced is useful for the proofs of correctness. It however does not provide one with a way to actually compute the relation. That is why we now provide an alternative definition that yields the same results but is more algorithmic in nature.

Definition 0.1.3. Let $\mathcal{A} = (Q, \Sigma, q_0, \delta, c)$ be a DPA and let $R \subseteq Q \times Q$ be a congruence relation. For each equivalence class λ of R, we define the *path refinement automaton* $\mathcal{G}_{\mathrm{PR}}^{R,\lambda}(p, q) = (Q_{\mathrm{PR}}, \Sigma, q_{0,\mathrm{PR}}, \delta_{\mathrm{PR}}, F_{\mathrm{PR}})$, which is a DFA.

- $Q_{\mathrm{PR}} = (Q \times Q \times c(Q) \times \{<, >, =\}) \cup \{\bot\}$
- $q_{0,\mathrm{PR}} = (p, q, \eta_k(c(p), c(q), \checkmark), \eta_x(c(p), c(q), \checkmark, =))$
- $\delta_{\mathrm{PR}}((p, q, k, x), a) = \begin{cases} (p', q', \eta_k(c(p'), c(q'), k), \eta_x(c(p'), c(q'), k, x)) & \text{if } p' \notin \lambda \\ (p', q', \eta_l(c(p'), c(q'), \checkmark), \eta_x(c(p'), c(q'), \checkmark, =)) & \text{if } p' \in \lambda \text{ and } (x = =) \\ \bot & \text{else} \end{cases}$

  where $p' = \delta(p, a)$ and $q' = \delta(q, a)$.
  $\eta_k(k_p, k_q, k) = \min_{\leq_\checkmark}\{k_p, k_q, k\}$
  $\eta_x(k_p, k_q, k, x) = \begin{cases} < & \text{if } k_p < k_q \text{ and } k_p <_\checkmark k \\ > & \text{if } k_q < k_p \text{ and } k_q <_\checkmark k \\ x & \text{if } k \leq_\checkmark k_p \text{ and } k \leq_\checkmark k_q \text{ and } (k \neq k_p \text{ or } k \neq k_q) \\ = & \text{else} \end{cases}$
- $F_{\mathrm{PR}} = Q_{\mathrm{PR}} \setminus \{\bot\}$

Theorem 0.1.3. *Let $\mathcal{A}$ be a DPA with a congruence relation R. Let λ be an equivalence class of R and $p, q \in \lambda$. Then $p \equiv_{\mathrm{PR}}^{R} q$ iff $L(\mathcal{G}_{\mathrm{PR}}^{R,\lambda}(p, q)) = \Sigma^*$.*

Proof. **If** Let $p \not\equiv_{\mathrm{PR}}^{R} q$. There is a word $w \in L_{\lambda \to \lambda}$ such that $(\delta^*(p, w), \delta^*(q, w)) \notin R$ or $\min\{c(\delta^*(p, u)) \mid u \sqsubset w\} \neq \min\{c(\delta^*(q, u)) \mid u \sqsubset w\}$.

In the second case, assume $\min\{c(\delta^*(p, u)) \mid u \sqsubset w\} < \min\{c(\delta^*(q, u)) \mid u \sqsubset w\}$ without loss of generality. Let $u \sqsubset w$ be the prefix of w with $\min\{c(\delta^*(p, u)) \mid u \sqsubset w\} = c(\delta^*(p, u))$.

Only If Let $L(\mathcal{G}_{\mathrm{PR}}^{R,\lambda}(p, q)) \neq \Sigma^*$. □

The differences between different $\mathcal{G}_{\mathrm{PR}}^{R,\lambda}$ for different λ are minor and the question whether the accepted language is universal boils down to a simple question of reachability. Thus, $\equiv_{\mathrm{PR}}^{R}$ can be computed in $\mathcal{O}(|\mathcal{G}_{\mathrm{PR}}^{R,\lambda}|)$ which is $\mathcal{O}(|Q|^2 \cdot |c(Q)|)$.